



COMP 412
FALL 2010

Code Optimization, Part II

Regional Techniques

Comp 412

Copyright 2010, Keith D. Cooper & Linda Torczon, all rights reserved.

Students enrolled in Comp 412 at Rice University have explicit permission to make copies of these materials for their personal use.

Faculty from other educational institutions may use these materials for nonprofit educational purposes, provided this copyright notice is preserved.



Last Lecture

Introduced concept of a redundant expression

An expression, $x+y$, is redundant at point p if, along each path from the procedure's entry point to p , $x+y$ has already been evaluated and neither x nor y has been redefined.

- If $x+y$ is redundant at p , we can save the results of those earlier evaluations and reuse them at p , avoiding evaluation
- In a single block, we need only consider one such path
- We developed an algorithm for redundancy elimination in a single basic block
- Two pieces to the problem
 - Proving that $x+y$ is redundant
 - Rewriting the code to eliminate the redundant evaluation
- Value numbering does both for straightline code

Local Value Numbering

(Recap)



The LVN Algorithm, with bells & whistles

for $i \leftarrow 0$ to $n-1$

1. get the value numbers V_1 and V_2 for L_i and R_i
2. if L_i and R_i are both constant then
evaluate $L_i \text{ Op}_i R_i$, assign it to T_i , and mark T_i as a constant
3. if $L_i \text{ Op}_i R_i$ matches an identity then
replace it with a copy operation or an assignment
4. if Op_i commutes and $V_1 > V_2$ then
swap V_1 and V_2
5. construct a hash key $\langle V_1, \text{Op}_i, V_2 \rangle$
6. if the hash key is already present in the table then
replace operation I with a copy into T_i and mark T_i with the VN
else
insert a new VN into table for hash key & mark T_i with the VN

Block is a sequence of n operations of the form

$$T_i \leftarrow L_i \text{ Op}_i R_i$$

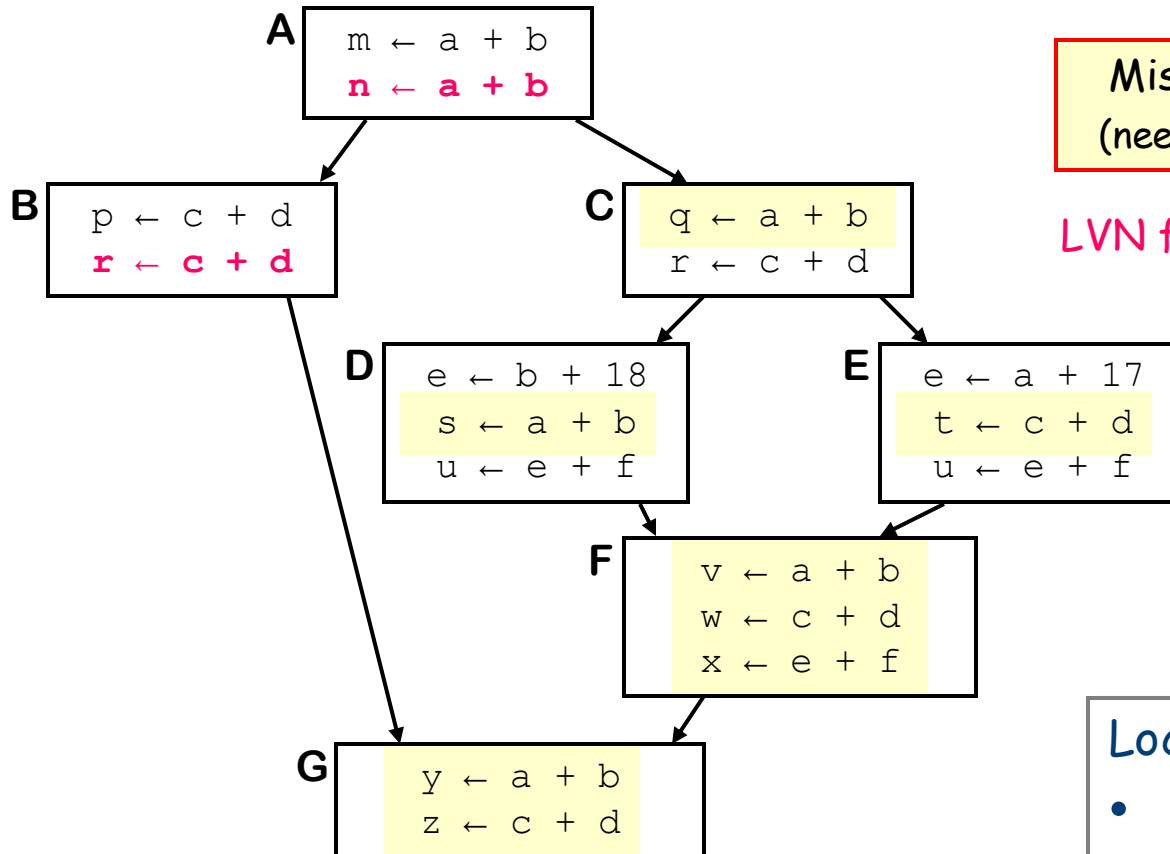
Constant folding

Algebraic identities

Commutativity



Local Value Numbering



Missed opportunities
(need stronger methods)

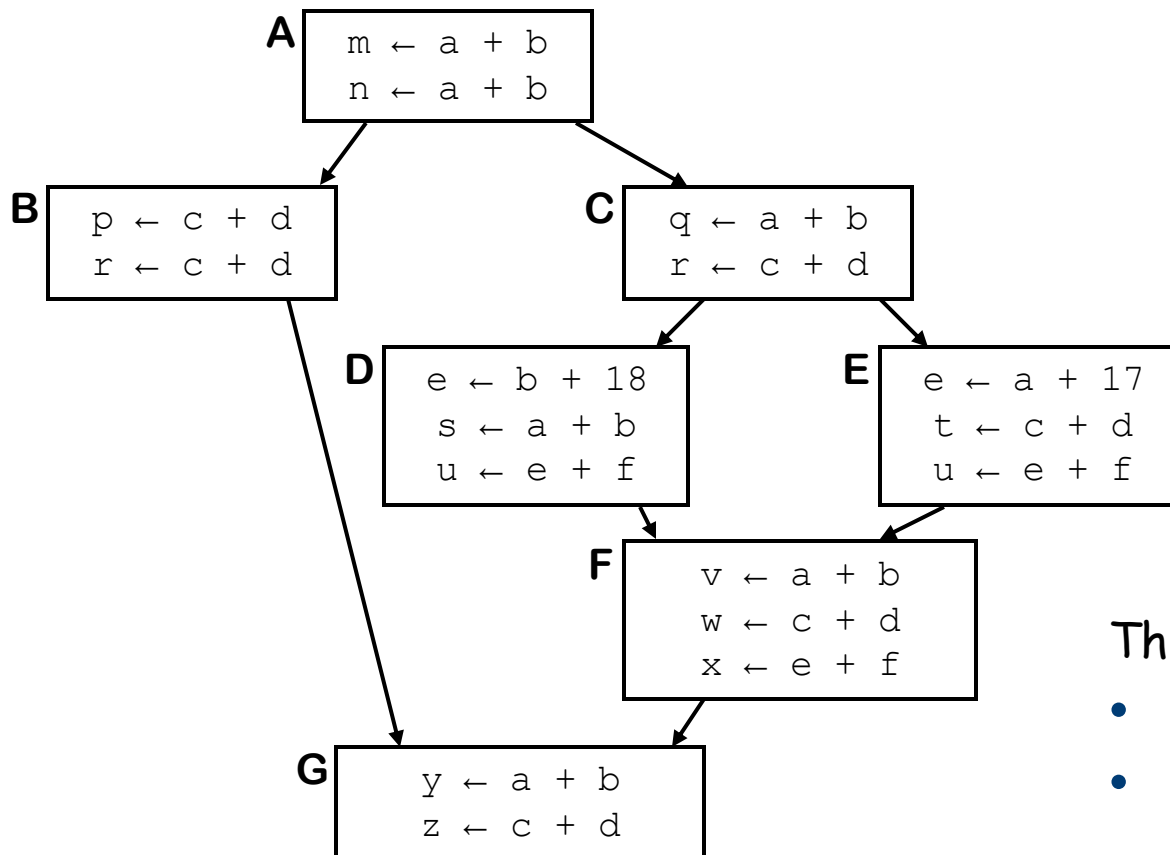
LVN finds these redundant ops

Local Value Numbering

- 1 block at a time
- Strong local results
- No cross-block effects



Terminology



Control-flow graph (CFG)

- Nodes for basic blocks
- Edges for branches
- Basis for much of program analysis & transformation

This CFG, $G = (N, E)$

- $N = \{A, B, C, D, E, F, G\}$
- $E = \{(A, B), (A, C), (B, G), (C, D), (C, E), (D, F), (E, F), (F, G)\}$
- $|N| = 7, |E| = 8$



Scope of Optimization

In scanning and parsing, "scope" refers to a region of the code that corresponds to a distinct name space.

In optimization "scope" refers to a region of the code that is subject to analysis and transformation.

- Notions are somewhat related
- Connection is not necessarily intuitive

Different scopes introduces different challenges & different opportunities

Historically, optimization has been performed at several distinct scopes.



Scope of Optimization

A basic block is a maximal length sequence of straightline code.

Local optimization

- Operates entirely within a single basic block
- Properties of block lead to strong optimizations

Regional optimization

- Operate on a region in the CFG that contains multiple blocks
- Loops, trees, paths, extended basic blocks

Whole procedure optimization (*intraprocedural*)

- Operate on entire CFG for a procedure
- Presence of cyclic paths forces analysis then transformation

Whole program optimization (*interprocedural*)

- Operate on some or all of the call graph (*multiple procedures*)
- Must contend with call/return & parameter binding



A Comp 412 Fairy Tale

We would like to believe optimization developed in an orderly fashion

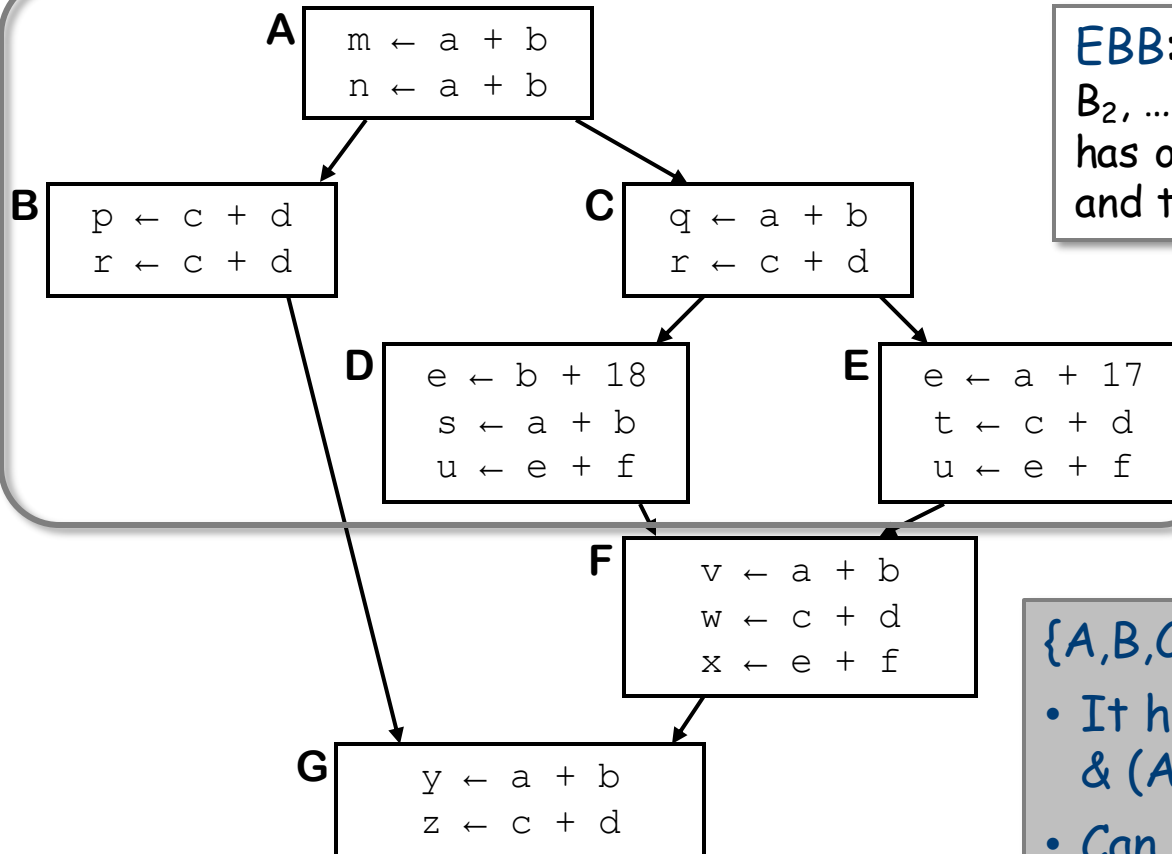
- Local methods led to regional methods
- Regional methods led to global methods
- Global methods led to interprocedural methods

It did not happen that way

- First compiler, FORTRAN, used both local & global methods
- Development has been scattershot & concurrent
- Scope appears to relate to the inefficiency being attacked, rather than the refinement of the inventor.



Superlocal Value Numbering



EBB: A maximal set of blocks B_1, B_2, \dots, B_n where each B_i , except B_1 , has only exactly one predecessor and that block is in the EBB.

$\{A, B, C, D, E\}$ is an EBB

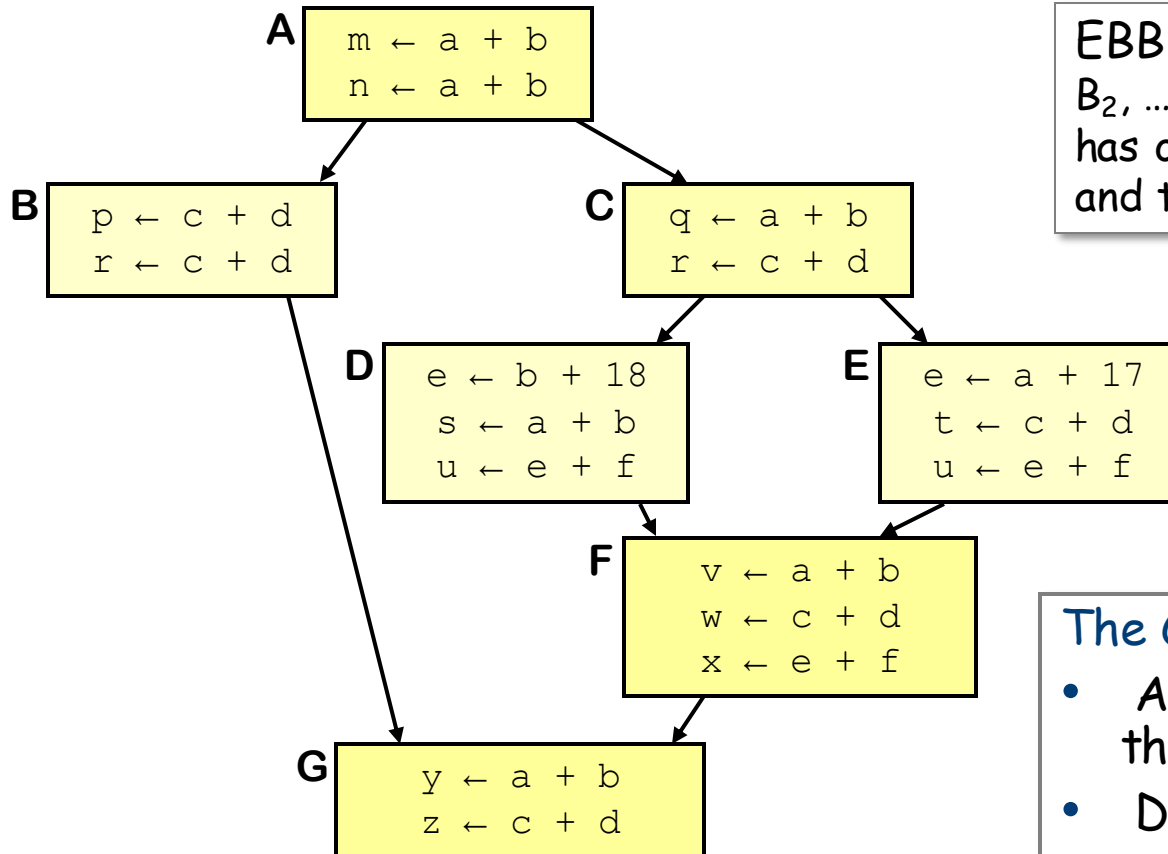
- It has 3 paths: (A, B) , (A, C, D) , & (A, C, E)
- Can sometimes treat each path as if it were a block

$\{F\}$ & $\{G\}$ are degenerate EBBs

Superlocal: "applied to an EBB"



Superlocal Value Numbering



EBB: A maximal set of blocks B_1, B_2, \dots, B_n where each B_i , except B_1 , has only exactly one predecessor and that block is in the EBB.

The Concept

- Apply local method to paths through the EBBs
- Do $\{A, B\}$, $\{A, C, D\}$, & $\{A, C, E\}$
- Obtain reuse from ancestors
- Avoid re-analyzing A & C
- Does not help with F or G

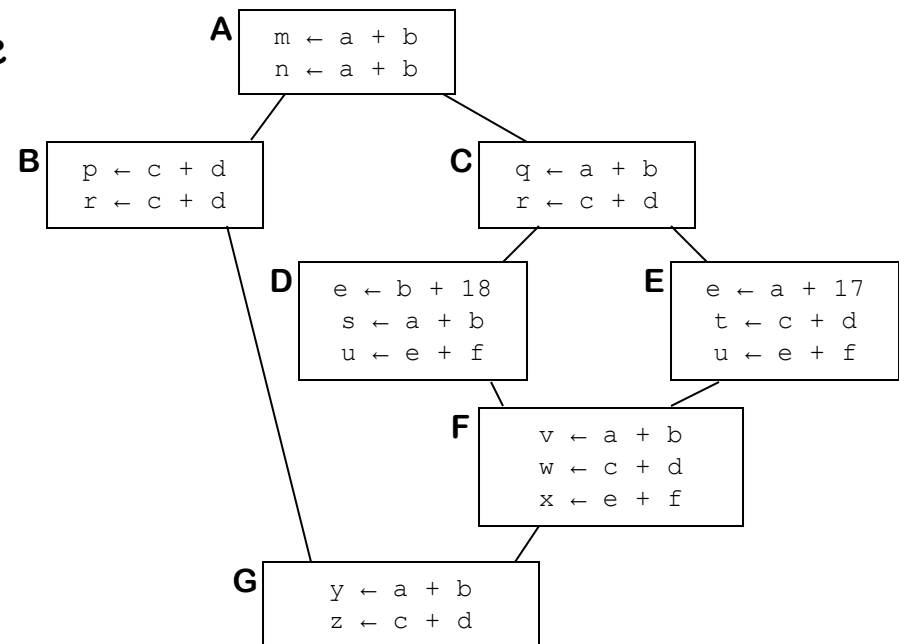


Superlocal Value Numbering

Efficiency

- Use A's table to initialize tables for B & C
- To avoid duplication, use a scoped hash table
 - A, AB, A, AC, ACD, AC, ACE, F, G
- Need a VN \rightarrow name mapping to handle kills
 - Must restore map with scope
 - Adds complication, not cost

"kill" is a re-definition of some name





Superlocal Value Numbering

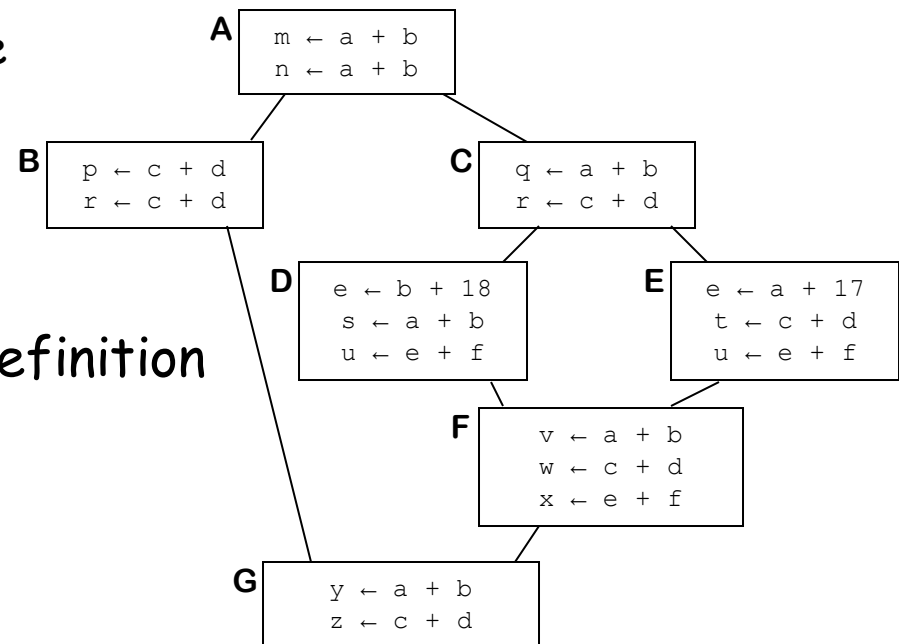
Efficiency

- Use A's table to initialize tables for B & C
- To avoid duplication, use a scoped hash table
 - A, AB, A, AC, ACD, AC, ACE, F, G
- Need a VN \rightarrow name mapping to handle kills
 - Must restore map with scope
 - Adds complication, not cost

"kill" is a re-definition of some name

To simplify matters

- Need *unique name* for each definition
- Makes name \rightarrow VN
- Use the SSA name space



The subscripted names from the earlier example are an instance of the SSA name space.

SSA Name Space

(locally)



Example (from earlier):

Original Code

$a_0 \leftarrow x_0 + y_0$
★ $b_0 \leftarrow x_0 + y_0$
 $a_1 \leftarrow 17$
★ $c_0 \leftarrow x_0 + y_0$

With VNs

$a_0^3 \leftarrow x_0^1 + y_0^2$
★ $b_0^3 \leftarrow x_0^1 + y_0^2$
 $a_1^4 \leftarrow 17$
★ $c_0^3 \leftarrow x_0^1 + y_0^2$

Rewritten

$a_0^3 \leftarrow x_0^1 + y_0^2$
★ $b_0^3 \leftarrow a_0^3$
 $a_1^4 \leftarrow 17$
★ $c_0^3 \leftarrow a_0^3$

Renaming:

- Give each value a unique name
- Makes it clear

Notation:

- While complex, the meaning is clear

Result:

- a_0^3 is available
- Rewriting just works



SSA Name Space

(in general)

Two principles

- Each name is defined by exactly one operation
- Each operand refers to exactly one definition

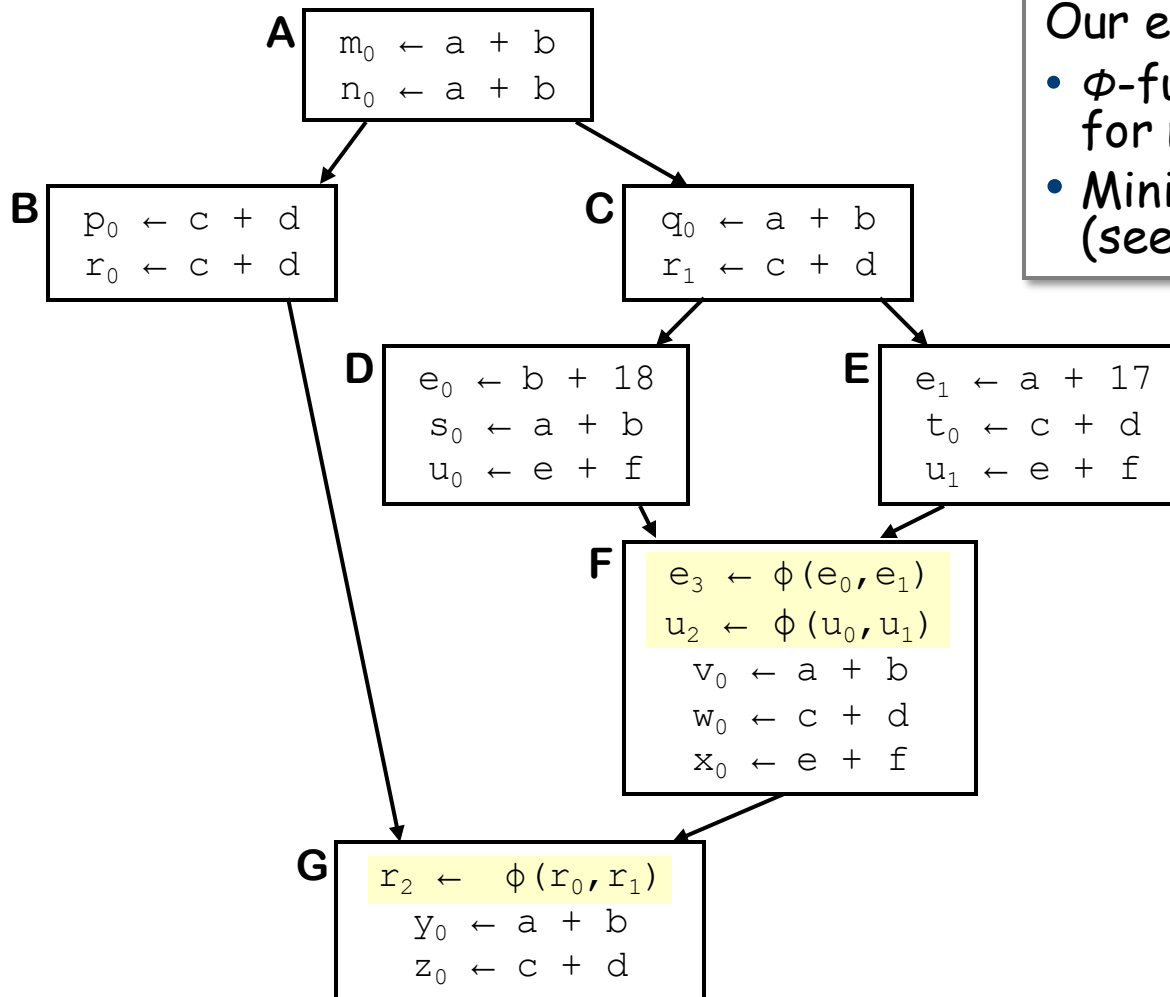
To reconcile these principles with real code

- Insert ϕ -functions at merge points to reconcile name space
- Add subscripts to variable names for uniqueness





Superlocal Value Numbering



Our example in SSA form

- ϕ -functions at join points for names that need them
- Minimal set of ϕ -functions (see Chapter 9 in EaC)



Superlocal Value Numbering

The SVN Algorithm

WorkList \leftarrow { entry block }

Blocks to process

Empty \leftarrow new table

Table for base case

while (*WorkList* is not empty)

remove a block b from WorkList

SVN(b, Empty)

SVN(Block, Table)

t \leftarrow new table for *Block*, with *Table* linked as surrounding scope

LVN(Block, t)

Use LVN for the work

for each successor s of Block

if s has just 1 predecessor

In the same EBB

then SVN(s, t)

else if s has not been processed

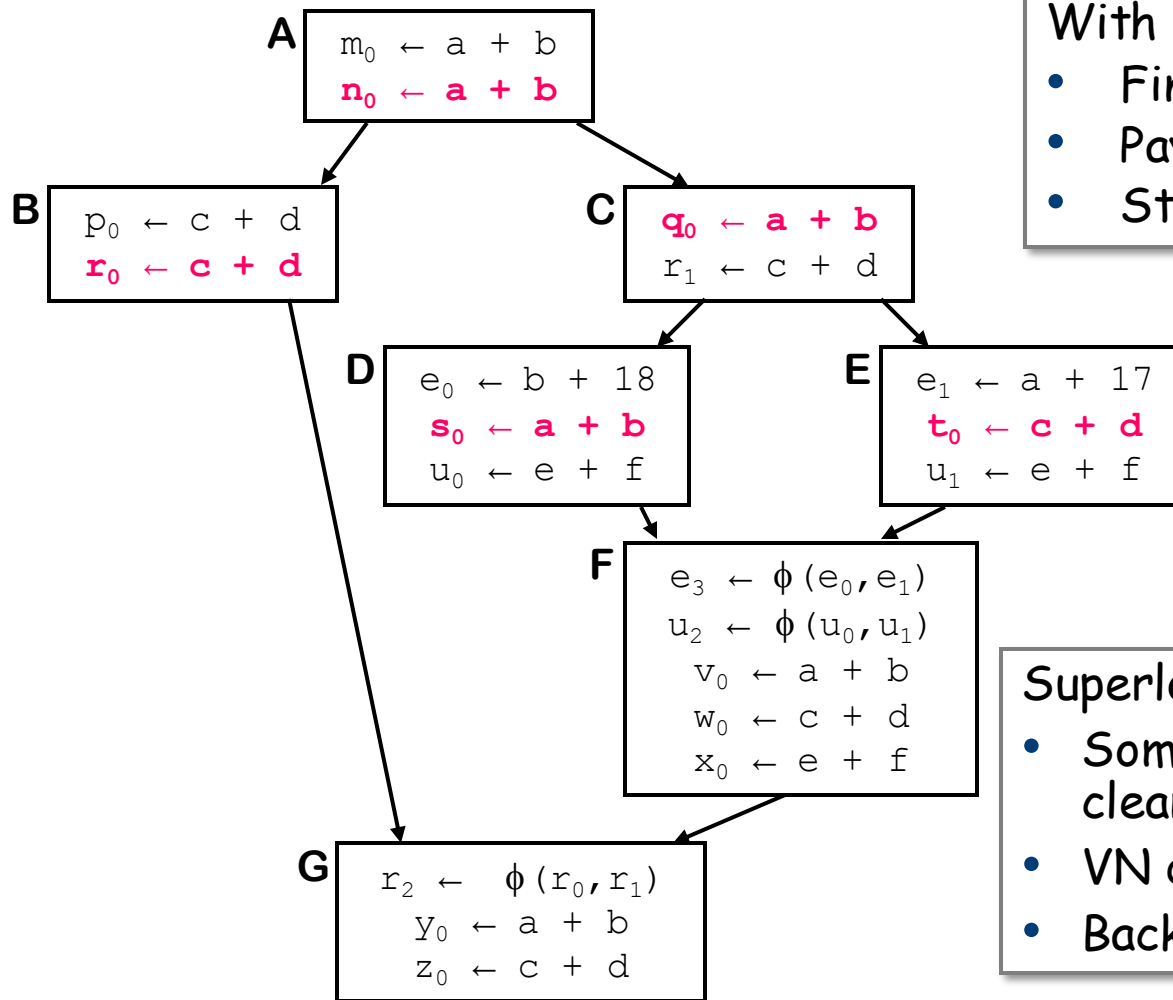
Starts a new EBB

then add s to WorkList

deallocate t



Superlocal Value Numbering



With all the bells & whistles

- Find more **redundancy**
- Pay minimal extra cost
- Still does nothing for F & G

Superlocal techniques


- Some local methods extend cleanly to superlocal scopes
- VN does not back up
- Backward motion causes probs



Loop Unrolling

Applications spend a lot of time in loops

- We can reduce loop overhead by unrolling the loop

| | | |
|--|--|---|
| <pre>do i = 1 to 100 by 1 a(i) ← b(i) * c(i) end</pre> |  | <pre>a(1) ← b(1) * c(1) a(2) ← b(2) * c(2) a(2) ← b(3) * c(3) ... a(100) ← b(100) * c(100)</pre> |
| | Complete unrolling | |

- Eliminated additions, tests, and branches
 - *Can subject resulting code to strong local optimization!*
- Only works with fixed loop bounds & few iterations
- The principle, however, is sound
- Unrolling is always safe, as long as we get the bounds right



Loop Unrolling

Unrolling by smaller factors can achieve much of the benefit

Example: unroll by 4

```
do i = 1 to 100 by 1  
  a(i) ← b(i) * c(i)  
end
```



```
do i = 1 to 100 by 4  
  a(i)    ← b(i) * c(i)  
  a(i+1) ← b(i+1) * c(i+1)  
  a(i+2) ← b(i+2) * c(i+2)  
  a(i+3) ← b(i+3) * c(i+3)  
end
```

Achieves much of the savings with lower code growth

- Reduces tests & branches by 25%
- LVN will eliminate duplicate adds and redundant expressions
- Less overhead per useful operation

But, it relied on knowledge of the loop bounds...



Loop Unrolling

Unrolling with unknown bounds

Need to generate guard loops

```
do i = 1 to n by 1
  a(i) ← b(i) * c(i)
end
```



Unroll by 4

```
i ← 1
do while (i+3 < n)
  a(i)    ← b(i) * c(i)
  a(i+1) ← b(i+1) * c(i+1)
  a(i+2) ← b(i+2) * c(i+2)
  a(i+3) ← b(i+3) * c(i+3)
  i ← i + 4
end
```

```
do while (i < n)
  a(i)    ← b(i) * c(i)
  i ← i + 1
end
```

Achieves most of the savings

- Reduces tests & branches by 25%
- LVN still works on loop body
- Guard loop takes some space

Can generalize to arbitrary upper & lower bounds, unroll factors



Loop Unrolling

One other unrolling trick

Eliminate copies at the end of a loop

```
t1 ← b(0)
do i = 1 to 100 by 1
  t2 ← b(i)
  a(i) ← a(i) + t1 + t2
  t1 ← t2
end
```



```
t1 ← b(0)
do i = 1 to 100 by 2
  t2 ← b(i)
  a(i) ← a(i) + t1 + t2
  t1 ← b(i+1)
  a(i+1) ← a(i+1) + t2 + t1
end
```

Unroll by LCM of copy-cycle lengths

- Eliminates the copies, which were a naming artifact
- Achieves some of the benefits of unrolling
 - Lower overhead, longer blocks for local optimization
- Situation occurs in more cases than you might suspect